

Class-level PGD-20 robustness remains uneven across adversarial-training recipes

